Getting Started with Java DB

Version 10.10

Derby Document build:
April 7, 2014, 12:27:27 PM (EDT)

Contents

Copyright....3

License....4

Relationship between Java DB and Apache Derby....8

Introduction to Derby....9
- **Deployment options**....9
- **System requirements**....9
- **Product documentation for Java DB**....10

Installing and configuring Derby....11
- **Installing Java DB**....11
- **Setting up your environment**....11
  - Choosing a method to run the Derby tools and startup utilities....11
  - Setting the environment variables....12
  - Syntax for the derbyrun.jar file....13
  - Manually setting the CLASSPATH environment variable....14
- **Using the Derby tools and startup utilities**....16
  - Running sysinfo....17
  - Running ij....18
  - Running dblook....19
- **Verifying the Derby system configuration**....20

Self-study tutorial for users new to Derby....21
- **Tutorial overview**....21
- **Activity 1: Run SQL using the embedded driver**....22
  - Creating a directory and copying scripts into the directory....22
  - Creating a Derby database and running SQL statements....23
- **Activity 2: Run SQL using the client driver**....25
- **Activity 3: Run a JDBC program using the embedded driver**....27
  - The WwdEmbedded program....28
- **Activity 4: Create and run a JDBC program using the client driver and Network Server**....32
- **What next with Derby?**....35

Quick start guide for experienced JDBC users....37
- **Environments in which Derby can run**....37
  - Embedded environment....37
  - Client/server environment....37
- **Available drivers**....37
- **Database connection URL**....37

Documentation conventions....39
- **Terminology**....39
- **SQL syntax**....39
- **Typographical conventions**....39

Derby libraries and scripts: complete reference....41
- **Libraries provided by Derby**....41
- **Scripts included with Derby**....42

Trademarks....44

Apache Software FoundationGetting Started with Java DBApache Derby

Copyright

Copyright 2004-2014 The Apache Software Foundation

Copyright 2005, 2014, Oracle and/or its affiliates. All rights reserved.

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at http://www.apache.org/licenses/LICENSE-2.0.

Related information

License

License

The Apache License, Version 2.0


```
                                Apache License
                           Version 2.0, January 2004
                        http://www.apache.org/licenses/

   TERMS AND CONDITIONS FOR USE, REPRODUCTION, AND DISTRIBUTION

   1. Definitions.

      "License" shall mean the terms and conditions for use,
      reproduction, and distribution as defined by Sections 1 through
      9 of this document.

      "Licensor" shall mean the copyright owner or entity authorized
      by the copyright owner that is granting the License.

      "Legal Entity" shall mean the union of the acting entity and all
      other entities that control, are controlled by, or are under
      common control with that entity. For the purposes of this
      definition, "control" means (i) the power, direct or indirect,
      to cause the direction or management of such entity, whether by
      contract or otherwise, or (ii) ownership of fifty percent (50%)
      or more of the outstanding shares, or (iii) beneficial ownership
      of such entity.

      "You" (or "Your") shall mean an individual or Legal Entity
      exercising permissions granted by this License.

      "Source" form shall mean the preferred form for making
      modifications, including but not limited to software source code,
      documentation source, and configuration files.

      "Object" form shall mean any form resulting from mechanical
      transformation or translation of a Source form, including but
      not limited to compiled object code, generated documentation,
      and conversions to other media types.

      "Work" shall mean the work of authorship, whether in Source or
      Object form, made available under the License, as indicated by a
      copyright notice that is included in or attached to the work
      (an example is provided in the Appendix below).

      "Derivative Works" shall mean any work, whether in Source or
      Object form, that is based on (or derived from) the Work and
      for which the editorial revisions, annotations, elaborations,
      or other modifications represent, as a whole, an original work
      of authorship.  For the purposes of this License, Derivative
      Works shall not include works that remain separable from, or
      merely link (or bind by name) to the interfaces of, the Work
      and Derivative Works thereof.

      "Contribution" shall mean any work of authorship, including
      the original version of the Work and any modifications or
      additions to that Work or Derivative Works thereof, that is
      intentionally submitted to Licensor for inclusion in the Work
      by the copyright owner or by an individual or Legal Entity
      authorized to submit on behalf of the copyright owner. For the
      purposes of this definition,
      "submitted" means any form of electronic, verbal, or written
      communication sent to the Licensor or its representatives,
      including but not limited to communication on electronic mailing
      lists, source code control systems, and issue tracking systems
```

```
    that are managed by, or on behalf of, the Licensor for the
    purpose of discussing and improving the Work, but excluding
    communication that is conspicuously marked or otherwise
    designated in writing by the copyright owner as "Not a
    Contribution."

    "Contributor" shall mean Licensor and any individual or Legal
    Entity on behalf of whom a Contribution has been received by
    Licensor and subsequently incorporated within the Work.

2. Grant of Copyright License. Subject to the terms and conditions
   of this License, each Contributor hereby grants to You a
   perpetual, worldwide, non-exclusive, no-charge, royalty-free,
   irrevocable copyright license to reproduce, prepare Derivative
   Works of, publicly display, publicly perform, sublicense, and
   distribute the Work and such Derivative Works in Source or
   Object form.

3. Grant of Patent License. Subject to the terms and conditions of
   this License, each Contributor hereby grants to You a perpetual,
   worldwide, non-exclusive, no-charge, royalty-free, irrevocable
   (except as stated in this section) patent license to make, have
   made, use, offer to sell, sell, import, and otherwise transfer
   the Work, where such license applies only to those patent claims
   licensable by such Contributor that are necessarily infringed by
   their Contribution(s) alone or by combination of their
   Contribution(s) with the Work to which such Contribution(s) was
   submitted. If You institute patent litigation against any entity
   (including a cross-claim or counterclaim in a lawsuit) alleging
   that the Work or a Contribution incorporated within the Work
   constitutes direct or contributory patent infringement, then any
   patent licenses granted to You under this License for that Work
   shall terminate as of the date such litigation is filed.

4. Redistribution. You may reproduce and distribute copies of the
   Work or Derivative Works thereof in any medium, with or without
   modifications, and in Source or Object form, provided that You
   meet the following conditions:

   (a) You must give any other recipients of the Work or
       Derivative Works a copy of this License; and

   (b) You must cause any modified files to carry prominent notices
       stating that You changed the files; and

   (c) You must retain, in the Source form of any Derivative Works
       that You distribute, all copyright, patent, trademark, and
       attribution notices from the Source form of the Work,
       excluding those notices that do not pertain to any part of
       the Derivative Works; and

   (d) If the Work includes a "NOTICE" text file as part of its
       distribution, then any Derivative Works that You distribute
       must include a readable copy of the attribution notices
       contained within such NOTICE file, excluding those notices
       that do not pertain to any part of the Derivative Works, in
       at least one of the following places: within a NOTICE text
       file distributed as part of the Derivative Works; within the
       Source form or documentation, if provided along with the
       Derivative Works; or, within a display generated by the
       Derivative Works, if and wherever such third-party notices
       normally appear. The contents of the NOTICE file are for
       informational purposes only and do not modify the License.
       You may add Your own attribution notices within Derivative
       Works that You distribute, alongside or as an addendum to
       the NOTICE text from the Work, provided that such additional
       attribution notices cannot be construed as modifying the
       License.

   You may add Your own copyright statement to Your modifications
```

```
   and may provide additional or different license terms and
   conditions for use, reproduction, or distribution of Your
   modifications, or for any such Derivative Works as a whole,
   provided Your use, reproduction, and distribution of the Work
   otherwise complies with the conditions stated in this License.

5. Submission of Contributions. Unless You explicitly state
   otherwise, any Contribution intentionally submitted for
   inclusion in the Work by You to the Licensor shall be under the
   terms and conditions of this License, without any additional
   terms or conditions.  Notwithstanding the above, nothing herein
   shall supersede or modify the terms of any separate license
   agreement you may have executed with Licensor regarding such
   Contributions.

6. Trademarks. This License does not grant permission to use the
   trade names, trademarks, service marks, or product names of the
   Licensor, except as required for reasonable and customary use
   in describing the origin of the Work and reproducing the content
   of the NOTICE file.

7. Disclaimer of Warranty. Unless required by applicable law or
   agreed to in writing, Licensor provides the Work (and each
   Contributor provides its Contributions) on an "AS IS" BASIS,
   WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
   implied, including, without limitation, any warranties or
   conditions of TITLE, NON-INFRINGEMENT, MERCHANTABILITY, or
   FITNESS FOR A PARTICULAR PURPOSE. You are solely responsible for
   determining the appropriateness of using or redistributing the
   Work and assume any risks associated with Your exercise of
   permissions under this License.

8. Limitation of Liability. In no event and under no legal theory,
   whether in tort (including negligence), contract, or otherwise,
   unless required by applicable law (such as deliberate and
   grossly negligent acts) or agreed to in writing, shall any
   Contributor be liable to You for damages, including any direct,
   indirect, special, incidental, or consequential damages of any
   character arising as a result of this License or out of the use
   or inability to use the Work (including but not limited to
   damages for loss of goodwill, work stoppage, computer failure or
   malfunction, or any and all other commercial damages or losses),
   even if such Contributor has been advised of the possibility of
   such damages.

9. Accepting Warranty or Additional Liability. While redistributing
   the Work or Derivative Works thereof, You may choose to offer,
   and charge a fee for, acceptance of support, warranty, indemnity,
   or other liability obligations and/or rights consistent with this
   License. However, in accepting such obligations, You may act only
   on Your own behalf and on Your sole responsibility, not on behalf
   of any other Contributor, and only if You agree to indemnify,
   defend, and hold each Contributor harmless for any liability
   incurred by, or claims asserted against, such Contributor by
   reason of your accepting any such warranty or additional
   liability.

END OF TERMS AND CONDITIONS

APPENDIX: How to apply the Apache License to your work.

   To apply the Apache License to your work, attach the following
   boilerplate notice, with the fields enclosed by brackets "[]"
   replaced with your own identifying information. (Don't include
   the brackets!)  The text should be enclosed in the appropriate
   comment syntax for the file format. We also recommend that a
   file or class name and description of purpose be included on the
   same "printed page" as the copyright notice for easier
   identification within third-party archives.
```

```
Copyright [yyyy] [name of copyright owner]

Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
implied.  See the License for the specific language governing
permissions and limitations under the License.
```

Relationship between Java DB and Apache Derby

Java DB is a relational database management system that is based on the Java programming language and SQL. Java DB is the Oracle release of the Apache Derby project, the Apache Software Foundation's (ASF) open source relational database project.

The Java DB product includes Derby without any modification whatsoever to the underlying source code.

Because Java DB and Derby have the same functionality, the Java DB documentation refers to the core functionality as Derby.

The Java DB 10.10 documentation is based on the Derby 10.10 documentation. References to "Derby" in the Java DB documentation should be understood as synonyms for "Java DB."

Oracle has made changes to the Apache Derby documentation. This manual is identical to *Getting Started with Derby*, with the following exceptions:

- Oracle has added this topic, "Relationship between Java DB and Apache Derby".
- In the titles of manuals, "Derby" has been changed to "Java DB".
- Path names that refer to locations within a Derby installation have been changed to refer to locations within a Java Development Kit (JDK) installation as needed. Specifically, demos and samples are in a different location within the JDK from their location in a Derby installation.
- In this manual, the topics "Product documentation for Java DB" and "Installing Java DB" have been changed slightly from the equivalent Derby topics.

Introduction to Derby

Welcome to Derby! Derby is a full-featured, open source relational database management system (RDBMS) that is based on Java technology and SQL.

Derby is written and implemented completely in the Java programming language. Derby provides users with a small-footprint standards-based database engine that can be tightly embedded into any Java based solution. Derby ensures data integrity and provides sophisticated transaction support. In the default configuration there is no separate database server to be installed or maintained by the end user.

The on-disk database format used by Derby is portable and platform-independent. You can move Derby databases from machine to machine without needing to modify the data. A Derby application can include a pre-built, populated database if it needs to, and that database will work in any Derby configuration.

For more information about Derby, visit the Derby website at http://db.apache.org/derby. The website provides pointers to the Derby Wiki and other resources, such as the derby-users mailing list, where you can ask questions about issues not covered in the documentation.

Before you install Derby, you should understand the deployment options and system requirements.

Deployment options

The Derby software distribution provides two basic deployment options (also referred to as frameworks), the simple embedded option and the Derby Network Server option.

Embedded

Refers to Derby being started by a simple single-user Java application. With this option Derby runs in the same Java virtual machine (JVM) as the application. Derby can be almost invisible to the end user because it is started and stopped by the application and often requires no administration.

The Derby documentation often refers to this as the embedded configuration or embedded mode.

Server (or Server-based)

Refers to Derby being started by an application that provides multi-user connectivity to Derby databases across a network. With this option Derby runs in the Java virtual machine (JVM) that hosts the Server. Applications connect to the Server from different JVMs to access the database. The Derby Network Server is part of the Derby software distribution and provides this type of framework for Derby. Derby also works well with other, independently developed Server applications.

The Derby documentation often refers to this as the Network Server configuration or client/server configuration.

System requirements

Derby is a database engine written completely in the Java programming language. The database will run in any certified Java Virtual Machine (JVM).

You must have a Java Development Kit (JDK) version 5 or higher installed on your computer. The JDK is required to perform the activities in the Self-study tutorial for users new to Derby.

To check that the correct version of the JDK is installed, issue the `java -version` command.

Product documentation for Java DB

The Java DB product documentation consists of the Java DB manuals and the API reference.

Getting Started with Java DB
: Describes how to install and configure Java DB. Includes a self-study tutorial for users new to Java DB and a quick-start guide for experienced JDBC users. This guide introduces the `dblook`, `ij`, and `sysinfo` tools, and the libraries and scripts that are included with Java DB.

Java DB Developer's Guide
: Describes the functionality and features of Java DB common to all deployments, such as Java DB's JDBC and SQL specifics, deploying Java DB applications, working with properties, security, and other advanced features.

Java DB Reference Manual
: Documents the implementation of the SQL language in Java DB. This guide provides reference information about the JDBC and JTA implementations, keywords, system tables, properties, and *SQLExceptions* in Java DB.

Tuning Java DB
: This guide offers performance tips, an in-depth discussion of performance, and information about the Java DB optimizer.

Java DB Tools and Utilities Guide
: Describes how to use the the Java DB tools such as `dblook`, `ij`, and `sysinfo`.

Java DB Server and Administration Guide
: Part One of this guide discusses configuring servers, how to program client applications, and database administration. In addition, some systems might require administrative tasks such as backing up databases. These tasks are independent of any server framework but are unique to multi-user or large systems.

  Part Two of this guide discusses administrative tasks such as backing up and restoring databases, importing and exporting data, storing Java code in the database, and debugging deadlocks.

Java DB API Reference
: An API Reference that is automatically generated for all public Java DB classes. No reference is provided for the JDBC API, which is part of the Java Platform, Standard Edition. For more information about the classes in the API, see the *Java DB Reference Manual*.

You can access the Java DB manuals and API Reference from the Java DB Documentation page.

Installing and configuring Derby

This section will help you install and configure Derby.

After you complete the installation:

- Access the Self-study tutorial to get up and running with Derby. This tutorial demonstrates how to use Derby in the embedded and client/server configurations.
- If you are an experienced JDBC programmer, you should also see the Quick start guide for experienced JDBC users.

Installing Java DB

Java DB is installed automatically as part of the Java SE Development Kit (JDK).

To obtain the JDK, navigate your web browser to http://www.oracle.com/technetwork/java/javase/downloads/index.html and click the Download button for the version of the JDK you want to install. Follow the instructions on subsequent pages. To obtain the demos and sample programs for Java DB, download and install the Demos and Samples for the JDK in addition to the JDK itself.

Java DB is installed in the *db* directory of the JDK installation. This distribution contains scripts and libraries. The installation contains the following subdirectories:

- The *bin* subdirectory contains the scripts for executing utilities and setting up the environment.
- The *lib* subdirectory contains the Java DB jar files.

For a Java DB installation, set the DERBY_HOME environment variable, described in Setting the environment variables, to `JAVA_HOME/db`.

After you install the Demos and Samples for the JDK, you can find the Java DB demos and samples in `JAVA_HOME/demo/db`.

Setting up your environment

Derby includes several scripts that run the Derby tools and utilities. How you run those scripts depends on the configuration that you want to use to access the Derby database and tools.

To set up your environment:

1. Choose the method that you want to use to run the Derby tools and utilities.
2. Set the environment variables.

Choosing a method to run the Derby tools and startup utilities

There are several ways that you can run the Derby tools and startup utilities.

The method that you choose to run the tools and utilities determines the environment variables that you set for Derby.

1. Choose the method that you want to use. The methods are listed in the following table.

   Table 1. Methods of running the Derby tools and utilities

Method	When to Use	Requirements
Run the tools using the command shell scripts that are provided with Derby.	Use this method when you want to run the tools with the least amount of typing and have installed the full `bin` distribution of Derby.	• Set the DERBY_HOME environment variable • Include the location of the `java` executable file (usually `JAVA_HOME\bin`) in the PATH environment variable • Include the `DERBY_HOME\bin` directory in the PATH environment variable
Run the tools using the `derbyrun.jar` archive file.	Use this method when: • You have only the Derby jar files available (see Libraries provided by Derby), and do not have the full `bin` distribution of Derby • You do not want to use the scripts that are provided with Derby	• Set the DERBY_HOME environment variable • Include the location of the `java` executable file (usually `JAVA_HOME\bin`) in the PATH environment variable The `derbyrun.jar` file must be in the same folder as the other Derby .jar files. For more information, see Syntax for the derbyrun.jar file.
Run the tools using the complete `java` syntax, or use Derby in a Java program.	Use this method when: • You do not have the `derbyrun.jar` file or the Derby scripts installed • You want to learn the full syntax of each command and understand the details of setting up the execution environment • You are developing application programs with Derby	• Set the DERBY_HOME environment variable • Include the location of the `java` executable file (usually `JAVA_HOME\bin`) in the PATH environment variable • You must know the full package name for the Java class that supports the tool • The CLASSPATH environment variable must be set to include the required JAR files For details on setting the CLASSPATH, see Manually setting the CLASSPATH environment variable.

2. Based on the requirements of the method that you chose to run the tools, follow the instructions to set the environment variables.

Setting the environment variables

There are several environment variables that must be set depending on the method that you selected to run the Derby tools and startup utilities.

As mentioned in choosing a method to run the Derby tools and startup utilities, you must set the DERBY_HOME environment variable so that you can use the command examples that are presented in this manual. Adding the Derby scripts directory to your command execution PATH makes the scripts easier to use and enables you to use the script examples in this manual. The CLASSPATH environment variable must be set if you are using Derby in a Java program or executing the tools using the `java` command.
The steps below show you how to set the environment variables in a command window. The settings are only valid for that window. If you close the command window or open a new command window, you must set the environment variables again.
Tip: You can also set environment variables permanently. For example, on Windows you can use the Control Panel to permanently set the environment variables.

To set the environment variables:

1. Set the DERBY_HOME environment variable to the location where you extracted the Derby *bin* distribution.

   For example, if you installed Derby in the `/opt/Derby_10` directory on UNIX or in the `c:\Derby_10` directory on Windows, use one of the commands in the following table to set the DERBY_HOME environment variable.

   Table 2. Commands to set the DERBY_HOME environment variable

Operating System	Command
UNIX (Korn Shell)	`export DERBY_HOME=/opt/Derby_10`
Windows	`set DERBY_HOME=c:\Derby_10`

2. Be certain that the `java` executable file, version 5 or higher, is in your command execution PATH. Open a command window and run the `java -version` command.
3. Add the `DERBY_HOME/bin` directory to the PATH environment variable so that you can run the Derby scripts from any directory, as shown in the following table.

   Table 3. Commands to set the PATH environment variable

Operating System	Command
UNIX (Korn Shell)	`export PATH="$DERBY_HOME/bin:$PATH"`
Windows	`set PATH=%DERBY_HOME%\bin;%PATH%` If you use the Control Panel to update your system PATH, add `%DERBY_HOME%\bin` to the end of the PATH environment variable.

   Tip: When the DERBY_HOME environment variable is set and the underlying `/bin` directory is included in the PATH environment variable, you can use shortened commands to start the Derby tools. Otherwise, either you must be in the directory where the script that starts the Derby tool is located, or you must specify the full path to the location of the script when you want to start the tool.

For more information on the scripts included in the *bin* distribution, see Scripts included with Derby.

Syntax for the derbyrun.jar file

The `derbyrun.jar` file is a special JAR file that simplifies how you invoke the Derby tools and the Network Server.

With the `derbyrun.jar` file, you can run the Derby tools and utilities using shortened names and you do not need to set the java CLASSPATH environment variable. The `derbyrun.jar` file must be in the same folder as the other Derby JAR files.

Syntax

The following table shows the syntax for using `derbyrun.jar` for each of the Derby tools.

Table 4. Syntax for using `derbyrun.jar`

Operating System	Command

UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar ij [-p` *`propertiesfile`*`] [`*`sql_script`*`]`
	`java -jar $DERBY_HOME/lib/derbyrun.jar sysinfo [-cp ...] [-cp help]`
	`java -jar $DERBY_HOME/lib/derbyrun.jar dblook [`*`arg`*`]*` (or no arguments for usage)
	`java -jar $DERBY_HOME/lib/derbyrun.jar server [`*`arg`*`]*` (or no arguments for usage)

Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar ij [-p` *`propertiesfile`*`] [`*`sql_script`*`]`
	`java -jar %DERBY_HOME%\lib\derbyrun.jar sysinfo [-cp ...] [-cp help]`
	`java -jar %DERBY_HOME%\lib\derbyrun.jar dblook [`*`arg`*`]*` (or no arguments for usage)
	`java -jar %DERBY_HOME%\lib\derbyrun.jar server [`*`arg`*`]*` (or no arguments for usage)

To see this syntax reminder, run the command `java -jar derbyrun.jar` with no arguments.

Additional information

You cannot use the `-cp` argument or the CLASSPATH environment variable to set CLASSPATH variables when you are using the `-jar` argument to start the `ij` tool. If you want to run the `ij` tool with a custom classpath, you cannot use the `-jar` argument. Instead, you have to use the full class name to start the `ij` tool (`java org.apache.derby.tools.ij`).

See Manually setting the CLASSPATH environment variable for more information about setting up the classpath and running the `ij` tool.

There is no such limitation when you run the `sysinfo` tool. See Running sysinfo for more information on running the `sysinfo` tool.

The `server` argument is a shortcut for running the `NetworkServerControl` tool. For details on using this tool, see the *Java DB Server and Administration Guide*.

For more information on using `derbyrun.jar` to run the `ij`, `sysinfo`, and `dblook` tools, see the *Java DB Tools and Utilities Guide*.

Manually setting the CLASSPATH environment variable

You can set the CLASSPATH environment variable in the operating system either temporarily, permanently, or at run time when you start your Java application and the JVM.

The classpath is a list of the class libraries that are needed by the JVM and other Java applications to run your program. There are scripts that are included with Derby that can set up the classpath to run the Derby tools.

If you want to call the Derby tools directly using Java and not using the scripts, you must manually set the CLASSPATH environment variable.

In most development environments, it is best to temporarily set the CLASSPATH environment variable in the command line shell where you are entering commands.

Derby provides several scripts in the DERBY_HOME`/bin` directory to help you set your classpath quickly. These scripts are:

setEmbeddedCP
: Use the setEmbeddedCP script to set the classpath when the database engine is used in embedded mode. This script adds the `derby.jar` and `derbytools.jar` files to the classpath.

setNetworkServerCP
: Use the setNetworkServerCP script to set the classpath when you want to start the network server. This script adds the `derbynet.jar` file to the classpath.

setNetworkClientCP
: Use the setNetworkClientCP script to set the classpath when you want to access databases using the network client. This script adds the `derbyclient.jar` and `derbytools.jar` files to the classpath.

To set the classpath temporarily, run the script that is appropriate for your environment every time that you open a new command window.

Note that these scripts behave slightly differently on UNIX systems and on Windows systems. On Windows systems, running the script in your command shell will set the environment variables for your shell.

On UNIX systems, you need to use the "dot" or "source" command to ensure that the script is run in the calling shell's environment. Otherwise, when a script is run, it launches a new shell for that script. The CLASSPATH is set within that shell. Once the script is complete, that shell exits and you are returned to your shell. The CLASSPATH is changed only for the shell that the script was running in. The CLASSPATH in your shell is unchanged.

The UNIX shell scripts are known to run successfully in the Bash shell, and you may need to modify them slightly if you are using a different UNIX shell.

Here are examples of running the scripts, then displaying the classpath that the script sets:

- On UNIX, using the "dot" command to invoke the script, a sample session might be as follows:

```
sh-2.05b$ echo $CLASSPATH
sh-2.05b$ DERBY_HOME=/derby/db-derby-10.X.Y.0-bin
sh-2.05b$ . $DERBY_HOME/bin/setEmbeddedCP
sh-2.05b$ echo $CLASSPATH
/derby/db-derby-10.X.Y.0-bin/lib/derby.jar:/derby/db-derby-10.X.Y.0-
bin/lib/derbytools.jar:
sh-2.05b$ java org.apache.derby.tools.ij
ij version 10.X
ij> quit;
```

- On UNIX, using the "source" command to invoke the script, a sample session might be as follows:

```
-bash-2.05b$ echo $CLASSPATH
-bash-2.05b$ export DERBY_HOME=/derby/db-derby-10.X.Y.0-bin
-bash-2.05b$ source $DERBY_HOME/bin/setEmbeddedCP
```

- On Windows, a sample session might be as follows:

```
C:\derby\db-derby-10.X.Y.0-bin\bin>echo %CLASSPATH%
%CLASSPATH%
C:\derby\db-derby-10.X.Y.0-bin\bin>set
 DERBY_HOME=C:\derby\db-derby-10.X.Y.0-bin
C:\derby\db-derby-10.X.Y.0-bin\bin>setEmbeddedCP.bat
C:\derby\db-derby-10.X.Y.0-bin\bin>echo %CLASSPATH%
C:\derby\DB-DERBY-10.X.Y.0-bin\lib\derby.jar;C:\derby\DB-DERBY-
10.X.Y.0-bin\lib\derbytools.jar;
C:\derby\db-derby-10.X.Y.0-bin\bin>java org.apache.derby.tools.ij
ij version 10.X
ij> quit;
```

For more information on running the `ij` and `sysinfo` utilities, see the *Java DB Tools and Utilities Guide*

Using the Derby tools and startup utilities

There are several tools and utilities that you might want to use as you begin to work with Derby.

The tools that are included with Derby are `dblook`, `ij`, and `sysinfo`. To run these tools see:

- Running sysinfo
- Running ij
- Running dblook

Derby includes a set of scripts that you can use to start the Derby tools. The scripts are located in the DERBY_HOME`/bin` directory. When you run these scripts, your CLASSPATH environment variable is set if you have not already set it, and remains set as long as you are running the tool.

Most of the examples in this guide that show how to use the Derby scripts to launch the Derby tools assume that you are using the embedded mode of the Derby database engine. Use the instructions below to run the scripts with the Network Server.

Running the scripts with the Network Server

To run the scripts with the Network Server, use the following commands:

- For the `sysinfo` tool, issue the command

```
NetworkServerControl sysinfo
```

- For the `ij` tool, issue the command

```
set DERBY_OPTS=-Dij.protocol=jdbc:derby://localhost/
```

  and then start `ij` by issuing the command `ij`.
- For the `dblook` tool, call the script and specify the `-d` option and the full connection URL to the Network Server database. For example:

```
dblook -d 'jdbc:derby://localhost/myDB;user=usr'
```

Additional Derby utilities

In addition, there are Derby utilities that are system procedures that you can call by using the `ij` tool. For example, there are system procedures that you can use to import and export external files. Instructions on how to use these system procedures are included in the *Java DB Server and Administration Guide* and the *Java DB Reference Manual*.

Running sysinfo

The Derby `sysinfo` tool displays information about your Java environment and your version of Derby.

The `sysinfo` utility prints system information to a console.

Choose the method that you will use to run the `sysinfo` script. The methods are listed in the following table.

Table 5. Methods of running the `sysinfo` utility

Method	When to Use	Command
Run `sysinfo` as a standalone command.	Use this method if you are relatively new to the Java programming language and new to Derby.	Follow the steps in Setting the environment variables before you run the `sysinfo` utility using this method. To run the `sysinfo` script from the command line, use: `sysinfo` The `sysinfo` script sets the appropriate environment variables, including the CLASSPATH, and runs the `sysinfo` utility.

Run `sysinfo` using the jar file that is located in the directory where `sysinfo` resides.	Use this method if you are new to Derby, but are familiar with the Java programming language.	You must set the DERBY_HOME environment variable before you can run the `sysinfo` utility using this method.
		On UNIX, the command is:
		`java [options] -jar $DERBY_HOME/lib/derbyrun.jar sysinfo`
		On Windows, the command is:
		`java [options] -jar %DERBY_HOME%\lib\derbyrun.jar sysinfo`

Run `sysinfo` using the `java` command.	Use this method if you are familiar with both the Java programming language and Derby, and you have already added the location of the `java` executable file to your command execution PATH.	You must set your CLASSPATH. Use the steps specified in Manually setting the CLASSPATH environment variable. Then specify the class name in the `java` command. For example: `java org.apache.derby.tools.sysinfo`

See "sysinfo" in the *Java DB Tools and Utilities Guide* for more information about using the `sysinfo` utility.

Running ij

The Derby `ij` tool is a JDBC tool that you can use to run scripts or interactive queries against a Derby database.

- Choose the method that you will use to run the `ij` script. The methods are listed in the following table.

  Table 6. Methods of running the `ij` utility

Method	When to Use	Command
Run `ij` as a standalone command.	Use this method if you are relatively new to the Java programming language and new to Derby.	Follow the steps in Setting the environment variables before you run the `ij` tool using this method. To run the `ij` script from the command line, use: `ij` You must add the `DERBY_HOME/bin` directory to your PATH environment variable before you can run the `ij` tool. The `ij` script sets the appropriate environment variables, including the CLASSPATH, and starts the `ij` tool.

Run `ij` using the jar file that is located in the directory where `ij` resides.	Use this method if you are new to Derby, but are familiar with the Java programming language.	You must set the DERBY_HOME environment variable before you can run the `ij` tool using this method.
		On UNIX, the command is:
		`java -jar $DERBY_HOME/lib/derbyrun.jar ij`
		On Windows, the command is:
		`java -jar %DERBY_HOME%\lib\derbyrun.jar ij`

Run `ij` using the `java` command.	Use this method if you are familiar with both the Java programming language and Derby, and you have already added the location of the `java` executable file to your command execution PATH.	You must set your CLASSPATH. Use the steps specified in Manually setting the CLASSPATH environment variable. Then specify the class name in the `java` command. For example: `java org.apache.derby.tools.ij`

- When you are ready to leave the `ij` tool, type:

```
ij> exit;
```

See the *Java DB Tools and Utilities Guide* for more information about the `ij` tool.

Running dblook

The Derby `dblook` utility is a Data Definition Language (DDL) generation utility.

The `dblook` utility is a simple utility that dumps all or parts of the DDL of a user-specified database to either a console or a file. The generated DDL can then be used for such things as recreating all or parts of a database, viewing a subset of the objects in a database (for example, those objects that pertain to specific tables and schemas), or documenting the schema of a database.

Choose the method that you will use to run the `dblook` script. The methods are listed in the following table.

Table 7. Methods of running the `dblook` utility

Method	When to Use	Command
Run `dblook` as a standalone command.	Use this method if you are relatively new to the Java programming language and new to Derby.	Follow the steps in Setting the environment variables before you run the `dblook` utility using this method. To run the `dblook` script from the command line use: `dblook -d connectionURL [options]` The `dblook` script sets the appropriate environment variables, including the CLASSPATH, and runs the `dblook` utility.

Run `dblook` using the jar file that is located in the directory where `dblook` resides.	Use this method if you are new to Derby, but are familiar with the Java programming language.	You must set the DERBY_HOME environment variable before you can run the `dblook` utility using this method.
		On UNIX, the command is:
		`java [options] -jar $DERBY_HOME/lib/derbyrun.jar dblook -d connectionURL [options]`
		On Windows, the command is:
		`java [options] -jar %DERBY_HOME%\lib\derbyrun.jar dblook -d connectionURL [options]`

Run `dblook` using the `java` command.	Use this method if you are familiar with both the Java programming language and Derby, and you have already added the location of the `java` executable file to your command execution PATH.	You must set your CLASSPATH. Use the steps specified in Manually setting the CLASSPATH environment variable. Then specify the class name in the `java` command. For example: `java org.apache.derby.tools.dblook -d connectionURL [options]`

See Using the Derby tools and startup utilities for an example of the use of the *connectionURL*. See "dblook" and its subtopics in the *Java DB Tools and Utilities Guide* for more information about using the `dblook` utility.

Verifying the Derby system configuration

You should confirm that your computer has the correct software installed and that the environment variables are set before you start working with Derby.

To check the Derby system configuration:

1. Verify that the `java` executable file, version 5 or higher, is in your command execution PATH by opening a command window and running the `java -version` command.

   The output from the command looks something like this:

   ```
   java version "1.7.0_07"
   Java(TM) SE Runtime Environment (build 1.7.0_07-b10)
   Java HotSpot(TM) Client VM (build 23.3-b01, mixed mode, sharing)
   ```

   The output you see might be different from what is shown here because the `java -version` command outputs vendor-specific information. If the command produces an error or the version listed is not 5 or higher, install a JDK before continuing.
2. Verify that the DERBY_HOME environment variable is set and points to the root directory of the Derby 10.10 installation. Open a command window and run the appropriate command for your system.

   If you installed Derby in the /opt/Derby_10 directory on UNIX or the c:\Derby_10 directory on Windows, the command that you use and the expected output are shown in the following table.

   Table 8. Verifying the DERBY_HOME setting

Operating System	Command	Output
UNIX (Korn Shell)	`echo $DERBY_HOME`	`/opt/Derby_10`
Windows	`echo %DERBY_HOME%`	`C:\Derby_10`

   If you receive an error or do not see the expected output, ensure that you have set the `DERBY_HOME` environment variable, as described in Setting the environment variables.

Note: You can also use the `sysinfo` utility to verify that the environment variables are set correctly. If the environment variables are set correctly, the `sysinfo` command displays information about your Java virtual machine (JVM) and the version of Derby that you have installed.

Self-study tutorial for users new to Derby

To help you get up and running with Derby as quickly as possible, this self-study guide highlights some of the more important features of Derby.

This tutorial uses a series of activities designed to demonstrate the use of Derby in the embedded and client/server configurations. After performing these activities, you will find Derby to be an easy-to-use and fully functional relational database management system (RDBMS).

Tutorial skills and software requirements

To perform the tutorial activities, you do not need to have any prior knowledge of Java software, the Java Database Connectivity (JDBC) API, or SQL. Each activity provides the complete command syntax that you need to use for each operation. Instructions include the syntax for both a UNIX Korn shell and a Windows command prompt.

This tutorial demonstrates, but does not teach, the Java, JDBC and SQL code that is presented. If you want a deeper understanding of these topics, you will need to consult additional reference materials.

To perform the tutorial activities, you must have:

- A Java Development Kit (JDK) version 5 or higher installed on your computer
- If you are using JDK 7, the JDK Demos and Samples installed on your computer
- The binary (`bin`) installation of Apache Derby version 10.10 installed on your computer
- A basic knowledge of the computer command line interface:
  - How to start a command shell or window
  - How to navigate the file system hierarchy

System configuration

You should verify the system configuration before you perform the tutorial activities.

Problems and feedback on the tutorial activities

If you experience problems with any aspect of the tutorial activities, send an e-mail message to `derby-user@db.apache.org`. Someone from the Derby community will respond to your message.

The questions and feedback received will be used to make this tutorial more useful.

Tutorial overview

Each activity in this self-study tutorial highlights an important feature of Derby.

Here is what you can expect to learn in each activity:

Activity 1:
Use the Derby `ij` tool to load the Derby embedded driver and start the Derby database engine. Create the database `firstdb` and the table `FIRSTTABLE`. Use a few basic SQL statements to insert and select data. The Derby message log `derby.log` and the database directories and files are introduced.

Activity 2:
Use Derby within a client/server configuration. Start the Derby Network Server, which will embed the Derby engine. In a separate process, use the Derby `ij` tool to load the Derby client driver and connect to the Network Server. Create a database called `seconddb` and the table `SECONDTABLE`. Use a few basic SQL statements to insert and select data.

Activity 3:
Load the Derby database engine from a simple Java JDBC program. Use the embedded driver to create the database `jdbcDemoDB` and the `WISH_LIST` table. Populate the table with text entered from the keyboard, then view a list of the records in the table. Walk through the code to understand the basic structure of a JDBC program that accesses a Derby database. The `CLASSPATH` variable and connection URL attribute `;shutdown=true` are introduced.

Activity 4:
Modify the Java JDBC program to load the client driver and connect to the Derby Network Server. Compile the altered program and test that the program operates as it did in the previous activity.

Activity 1: Run SQL using the embedded driver

In this activity, you will use the embedded driver to load the Derby database engine, create and connect to a database, and use some basic SQL statements.

To run SQL using the embedded driver:

1. Create a directory and copy scripts into the directory
2. Create a Derby database and run SQL statements

Creating a directory and copying scripts into the directory

To help you complete this activity, you must create a directory and copy several scripts into the new directory. You will use these scripts to quickly add tables and data to a new Derby database.

Tip: A command prompt appears after each command is run. If an error occurs, verify the syntax and retype the command.

1. Open a command window and change to a directory where you want to store the files that you create during the self-study tutorial activities.
2. Create the `DERBYTUTOR` directory, as shown in the following table. `DERBYTUTOR` will be your working directory for this activity.

   Table 9. Command to create the `DERBYTUTOR` directory

Operating System	Command
UNIX (Korn Shell)	`mkdir DERBYTUTOR`
Windows	`md DERBYTUTOR`

3. Change to the `DERBYTUTOR` directory, as shown in the following table.

   Table 10. Command to change to the `DERBYTUTOR` directory

Operating System	Command
UNIX (Korn Shell)	`cd DERBYTUTOR`
Windows	`cd DERBYTUTOR`

4. Copy the SQL scripts from the JDK `demo\db\programs\toursdb` subdirectory into the `DERBYTUTOR` directory, as shown in the following table. You will use these scripts to create tables and add data to a new database, `toursdb`. See Installing Java DB for information on obtaining the demos.

   Table 11. Command to copy the SQL scripts

Operating System	Command
UNIX (Korn Shell)	`cp $JAVA_HOME/demo/db/programs/toursdb/*.sql .`
Windows	`copy %JAVA_HOME%\demo\db\programs\toursdb\*.sql .`

> Important: Include the dot (`.`) at the end of the command so that the files are copied to the correct location.

5. Verify that the files were copied to the `DERBYTUTOR` directory, as shown in the following table.

Table 12. Command to verify that the files were copied

Operating System	Command
UNIX (Korn Shell)	`ls`
Windows	`dir`

Creating a Derby database and running SQL statements

Now, you will use the Derby `ij` tool to load the Derby database engine. You will use the Derby embedded driver to create and connect to the `firstdb` database. You will also use a few basic SQL statements to create and populate a table.

1. Run the Derby `ij` tool.

   If you included the `DERBY_HOME/bin` directory in your PATH environment variable, type:

   ```
   ij
   ```

   Otherwise, you can use the `java` command to start the `ij` tool, as shown in the following table.

   Table 13. Command to start the `ij` tool

Operating System	Command
UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar ij` `ij version 10.10`
Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar ij` `ij version 10.10`

2. Create the database and open a connection to the database using the embedded driver.

   ```
   CONNECT 'jdbc:derby:firstdb;create=true';
   ```

   Description of connection command:

   `connect`
   The `ij` command to establish a connection to a database. The Derby connection URL is enclosed in single quotation marks. An `ij` command can be in either uppercase or lowercase.

   `jdbc:derby:`
   The JDBC protocol specification for the Derby driver.

   `firstdb`

The name of the database. The name can be any string. Because no filepath is specified, the database is created in the default working directory (`DERBYTUTOR`).

`;create=true`

The Derby *URL attribute* that is used to create a database. Derby does not have an SQL `create database` command.

`;`

The semicolon is the `ij` command terminator.

3. Create a table with two columns using standard SQL.

```
CREATE TABLE FIRSTTABLE
    (ID INT PRIMARY KEY,
    NAME VARCHAR(12));
0 rows inserted/updated/deleted
```

4. Insert three records.

```
INSERT INTO FIRSTTABLE VALUES
    (10,'TEN'),(20,'TWENTY'),(30,'THIRTY');
3 rows inserted/updated/deleted
```

5. Perform a simple select of all records in the table.

```
SELECT * FROM FIRSTTABLE;
ID         |NAME
------------------------
10         |TEN
20         |TWENTY
30         |THIRTY

3 rows selected
```

6. Perform a qualified select of the record with column ID=20.

```
SELECT * FROM FIRSTTABLE WHERE ID=20;
ID         |NAME
------------------------
20         |TWENTY

1 row selected
```

7. Optional: Create and populate additional tables and other schema objects.

   a. Load the SQL script `ToursDB_schema.sql`.

```
run 'ToursDB_schema.sql';

ij> ...
CREATE TABLE AIRLINES
   (
      AIRLINE CHAR(2) NOT NULL ,
      AIRLINE_FULL VARCHAR(24),
      BASIC_RATE DOUBLE PRECISION,
   ...
0 rows inserted/updated/deleted
   ... Other output messages not shown ...
```

   b. Populate the tables with data by running the script `loadTables.sql`.

```
run 'loadTables.sql';

ij> run 'loadCOUNTRIES.sql';
ij> insert into COUNTRIES values ( 'Afghanistan','AF','Asia');
1 row inserted/updated/deleted
ij> insert into COUNTRIES values ( 'Albania','AL','Europe');
1 row inserted/updated/deleted
   ... Other output messages not shown ...
```

8. Exit the `ij` tool.

```
exit;
```

You should be returned to the `DERBYTUTOR` directory.

9. Browse the most important files created by this activity:

   - The `derby.log` file. This file is a message and error log that, under normal circumstances, contains a set of startup messages and a shutdown message.

```
----------------------------------------------------------------
Thu Sep 13 09:52:15 EDT 2012:
Booting Derby version The Apache Software Foundation - Apache
  Derby - 10.10.0.0 - (1384314):
  instance a816c00e-0139-bfe6-bff8-000000a155b8
on database directory C:\DERBYTUTOR\firstdb
  with class loader sun.misc.Launcher$AppClassLoader@9931f5
Loaded from file:C:\db-derby-10.10.0.0-bin\lib\derby.jar
java.vendor=Oracle Corporation
java.runtime.version=1.7.0_07-b11
user.dir=C:\DERBYTUTOR
os.name=Windows XP
os.arch=x86
os.version=5.1
derby.system.home=C:\DERBYTUTOR
Database Class Loader started - derby.database.classpath=''
----------------------------------------------------------------
Thu Sep 13 09:53:21 EDT 2012: Shutting down Derby engine
----------------------------------------------------------------
Thu Sep 13 09:53:21 EDT 2012:
Shutting down instance a816c00e-0139-bfe6-bff8-000000a155b8 on
  database directory C:\DERBYTUTOR\firstdb
  with class loader sun.misc.Launcher$AppClassLoader@9931f5
----------------------------------------------------------------
```

   - The `firstdb` database directory. Within the directory are the subdirectories `seg0` (containing the data files) and `log` (containing the transaction log files).

Activity 2: Run SQL using the client driver

This activity uses Derby within a client/server configuration by using the Network Server. The `ij` tool is the client application that connects to the Derby Network Server. In this activity, you create a database called `seconddb` and run some basic SQL statements.

This activity assumes that you know how to open a command shell and change to the `DERBYTUTOR` directory.

You use two command windows (referred to as Shell-1 and Shell-2) in this activity. You use Shell-1 to start the Derby Network Server and display Network Server messages. You use Shell-2 to establish a client connection to the Network Server using `ij` and then perform some basic SQL operations.

1. Open a command window (Shell-1) and change to the `DERBYTUTOR` directory.
2. Start the Network Server, as shown in the following table.

Table 14. Command to start the Network Server

Operating System	Command

UNIX (Korn Shell)	**`java -jar $DERBY_HOME/lib/derbyrun.jar server start`**
	`Thu Sep 13 16:45:36 EDT 2012 : Security manager installed using the Basic server security policy.`
	`Thu Sep 13 16:45:37 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) started and ready to accept connections on port 1527`

Windows	

Operating System	Command

	`java -jar %DERBY_HOME%\lib\derbyrun.jar server start`
	`Thu Sep 13 16:45:36 EDT 2012 : Security manager installed using the Basic server security policy.`
	`Thu Sep 13 16:45:37 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) started and ready to accept connections on port 1527`

A Network Server startup message appears in the Shell-1 command window.

3. Open another command window (Shell-2). Change to the `DERBYTUTOR` directory.
4. Start `ij`.

   If you included the `DERBY_HOME/bin` directory in your PATH environment variable, type:

   ```
   ij
   ```

   Otherwise, you can use the `java` command to start the `ij` tool, as shown in the following table.

   Table 15. Command to start the `ij` tool

Operating System	Command
UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar ij` `ij version 10.10`
Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar ij` `ij version 10.10`

You will enter all subsequent commands from the network client, so you will type the commands in the Shell-2 command window.

5. Create and open a connection to the database using the client driver.

   ```
   CONNECT 'jdbc:derby://localhost:1527/seconddb;create=true';
   ```

   Remember: A client connection URL contains a hostname and a port number. For example: `//localhost:1527/`

6. Create a table with two columns (ID and NAME) using the following SQL statement:

   ```
   CREATE TABLE SECONDTABLE
       (ID INT PRIMARY KEY,
       NAME VARCHAR(14));
   0 rows inserted/updated/deleted
   ```

7. Insert three records into the table.

   ```
   INSERT INTO SECONDTABLE VALUES
       (100,'ONE HUNDRED'),(200,'TWO HUNDRED'),(300,'THREE HUNDRED');
   3 rows inserted/updated/deleted
   ```

8. Select all of the records in the table.

   ```
   SELECT * FROM SECONDTABLE;
   ID          |NAME
   ------------------------
   100         |ONE HUNDRED
   200         |TWO HUNDRED
   300         |THREE HUNDRED

   3 rows selected
   ```

9. Select a subset of records from the table by specifying a `WHERE` clause.

```
SELECT * FROM SECONDTABLE WHERE ID=200;
ID          |NAME
------------------------
200         |TWO HUNDRED

1 row selected
```

10. Exit `ij`.

```
exit;
```

11. Shut down the Derby Network Server, as shown in the following table.

Table 16. Command to shut down the Network Server

Operating System	Command
UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar server shutdown` `Thu Sep 13 16:50:37 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) shutdown`
Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar server shutdown` `Thu Sep 13 16:50:37 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) shutdown`

The server shutdown confirmation appears in both command windows.

Activity notes

The client connection URL contains network information (hostname and portnumber) not found in the URL for an embedded connection. This information tells the client driver the location of the Network Server. The client driver sends requests to and receives responses from the Network Server.

In this activity the Derby database engine is embedded in the Network Server and returns data to the `ij` client (a client/server configuration). In contrast, establishing a connection using an embedded URL (one without `//localhost:1527/`) would have caused the Derby engine to be embedded in the `ij` application (an embedded configuration).

In this configuration, multiple client programs can connect to the Network Server and access the database simultaneously. (This activity does not demonstrate this capability.)

Activity 3: Run a JDBC program using the embedded driver

This activity loads the Derby database engine using a simple Java JDBC program.

This activity assumes that you have opened a command window and navigated to the `DERBYTUTOR` directory.

JDBC is the Java Database Connectivity API and is also the native API for Derby. The program uses the embedded driver to create the `jdbcDemoDB` database (if the database does not exist) and then connect to the database. You can then populate a table within the database with text. The program demonstrates some basic JDBC processing along with related error handling.

The Java compiler and runtime use the classpath, specified by the `CLASSPATH` environment variable, to locate the binary files (jar files and class files) that are needed to run Derby and other Java applications. Before performing this activity, you need to set the classpath and compile the `WwdEmbedded.java` program.

1. Copy the program files into the `DERBYTUTOR` directory and set the `CLASSPATH` environment variable, as shown in the following table. See Installing Java DB for information on obtaining the demos.

   Table 17. Commands to copy the program files and set the `CLASSPATH` environment variable

Operating System	Command
UNIX (Korn Shell)	`cp $JAVA_HOME/demo/db/programs/workingwithderby/* .` `export CLASSPATH=$DERBY_HOME/lib/derby.jar:.`
Windows	`copy %JAVA_HOME%\demo\db\programs\workingwithderby\* .` `set CLASSPATH=%DERBY_HOME%\lib\derby.jar;.`

   > Important: Include the dot (`.`) at the end of each command so that your current working directory is included in the classpath and the files are copied to the correct location.
2. Compile the program source files.

   The sample program is contained in two source files: `WwdEmbedded.java` and `WwdUtils.java`. Issue the following command to compile both at the same time:

   ```
   javac WwdEmbedded.java WwdUtils.java
   ```

   > Important: A command prompt appears if the compilation is successful. The binary files `WwdEmbedded.class` and `WwdUtils.class` are created. If an error message appears, verify that the JDK is properly installed.
3. Run the program.

   The `WwdEmbedded.java` program populates a table with wish-list items. The program prompts you for text input (up to 32 characters), stores the text input in a database table, and then lists the items stored in the table. The program continues to ask for wish-list items until the you type the command `exit` or a problem is encountered. Some basic information on program progress is displayed at the beginning and the end of the program.

   ```
   java WwdEmbedded
   org.apache.derby.jdbc.EmbeddedDriver loaded.
   Connected to database jdbcDemoDB
    . . . . creating table WISH_LIST
   Enter wish-list item (enter exit to end):
   a peppermint stick
    ____________________________________________________
   On 2012-09-14 09:46:37.77 I wished for a peppermint stick
    ____________________________________________________
   Enter wish-list item (enter exit to end):
   a long vacation
    ____________________________________________________
   On 2012-09-14 09:46:37.77 I wished for a peppermint stick
   On 2012-09-14 09:46:51.654 I wished for a long vacation
    ____________________________________________________
   Enter wish-list item (enter exit to end):
   exit
   Closed connection
   Database shut down normally
   Getting Started With Derby JDBC program ending.
   ```

The WwdEmbedded program

This section describes the `WwdEmbedded.java` program, highlighting details specific to accessing a Derby database from a JDBC program.

Most of the code related to the database activities performed is included in this section, but you might find it helpful to open the program file and follow along in a text viewer or editor. The *SECTION NAMES* referred to in this section can be found in the comments within the program code and serve as cross-reference points between this section and the Java program. The program uses methods from the `WwdUtils` class. The utility class code is not described here but is available for review in the file `WwdUtils.java`.

Initialize the program

INITIALIZATION SECTION: The initial lines of code identify the Java packages used in the program, then set up the Java class `WwdEmbedded` and the `main` method signature. Refer to a Java programming guide for information on these program constructs.

```
import java.sql.*;

public class WwdEmbedded
{
    public static void main(String[] args)
   {
```

Define key variables and objects

DEFINE VARIABLES SECTION: The initial lines of the `main` method define the variables and objects used in the program. This example uses variables to store the information needed to connect to the Derby database. The use of variables for this information makes it easy to adapt the program to other configurations and other databases.

`driver`
: Stores the name of the Derby embedded driver.

`dbName`
: Stores the name of the database.

`connectionURL`
: Stores the Derby connection URL that is used to access the database.

`createString`
: Stores the SQL `CREATE` statement for the `WISH_LIST` table.

```
String driver = "org.apache.derby.jdbc.EmbeddedDriver";
String dbName="jdbcDemoDB";
String connectionURL = "jdbc:derby:" + dbName + ";create=true";
...
String createString = "CREATE TABLE WISH_LIST  "
  +  "(WISH_ID INT NOT NULL GENERATED ALWAYS AS IDENTITY "
  ...
  +  " WISH_ITEM VARCHAR(32) NOT NULL) " ;
```

Start the Derby engine

LOAD DRIVER SECTION: Loading the Derby embedded JDBC driver starts the Derby database engine. The `try` and `catch` block (the Java error-handling construct) catches the exceptions that may occur. A problem here is usually due to an incorrect classpath setting.

```
String driver = "org.apache.derby.jdbc.EmbeddedDriver";
...
try {
    Class.forName(driver);
} catch(java.lang.ClassNotFoundException e) {
  ...
}
```

Boot the database

BOOT DATABASE SECTION: The `DriverManager` class loads the database using the Derby connection URL stored in the variable `connectionURL`. This URL includes the parameter `;create=true` so that the database will be created if it does not already exist. The primary `try` and `catch` block begins here. This construct handles errors for the database access code.

```
String connectionURL = "jdbc:derby:" + dbName + ";create=true";
...
try {
    conn = DriverManager.getConnection(connectionURL);
    ...  <most of the program code is contained here>
}  catch (Throwable e)  {
   ...
}
```

Set up program to execute SQL

INITIAL SQL SECTION: The program initializes the objects needed to perform subsequent SQL operations and checks to see if the required data table exists.

The statement object `s` is initialized. If the utility method `WwdUtils.wwdChk4Table` does not find the `WISH_LIST` table, the statement object's `execute` method creates the table by executing the SQL stored in the variable `createString`.

```
s = conn.createStatement();
if (! WwdUtils.wwdChk4Table(conn))
{
   System.out.println (" . . . . creating table WISH_LIST");
   s.execute(createString);
}
```

The `INSERT` statement used to add data to the table is bound to the prepared statement object `psInsert`. The prepared statement uses the question mark parameter `?` to represent the data that will be inserted by the user. The program sets the actual value to be inserted later on, before executing the SQL. This is the most efficient way to execute SQL statements that will be used multiple times.

```
psInsert = conn.prepareStatement
   ("insert into WISH_LIST(WISH_ITEM) values (?)");
```

Interact with the database

ADD / DISPLAY RECORD SECTION: This section uses the utility method `WwdUtils.getWishItem` to gather information from the user. It then uses the objects set up previously to insert the data into the `WISH_LIST` table and then display all records. A standard `do` loop causes the program to repeat this series of steps until the user types `exit`. The data-related activities performed in this section are as follows:

- The `setString` method sets the substitution parameter of the `psInsert` object to the value typed by the user. Then the `executeUpdate` method performs the database insert.

  ```
  psInsert.setString(1,answer);
  psInsert.executeUpdate();
  ```

- The statement object `s` is used to select all the records in the `WISH_LIST` table and store them in the `ResultSet` named `myWishes`.

  ```
  myWishes = s.executeQuery("select ENTRY_DATE, WISH_ITEM
                  from WISH_LIST order by ENTRY_DATE");
  ```

  The `while` loop reads each record in turn by calling the `next` method. The `getTimestamp` and `getString` methods return specific fields in the record in the proper format. The fields are displayed using rudimentary formatting.

```
while (myWishes.next())
{
    System.out.println("On " + myWishes.getTimestamp(1) +
       " I wished for " + myWishes.getString(2));
}
```

Close the `ResultSet` to release the memory being used.

```
myWishes.close();
```

Shut down the database

DATABASE SHUTDOWN SECTION: If an application starts the Derby engine, the application should shut down all databases before exiting. The attribute `;shutdown=true` in the Derby connection URL performs the shutdown. When the Derby engine is shutdown, all booted databases will automatically shut down. The shutdown process cleans up records in the transaction log to ensure a faster startup the next time the database is booted.

Tip: You can shut down individual databases without shutting down the engine by including the database name in the connection URL.

"This section verifies that the embedded driver is being used, then issues the shutdown command and catches the shutdown exception to confirm that the Derby engine shut down cleanly. The shutdown status is displayed before the program exits.

```
if (driver.equals("org.apache.derby.jdbc.EmbeddedDriver")) {
   boolean gotSQLExc = false;
   try {
      DriverManager.getConnection("jdbc:derby:;shutdown=true");
   } catch (SQLException se)  {
      if ( se.getSQLState().equals("XJ015") ) {
         gotSQLExc = true;
      }
   }
   if (!gotSQLExc) {
      System.out.println("Database did not shut down normally");
   }  else  {
      System.out.println("Database shut down normally");
   }
}
```

> Important: The XJ015 error (successful shutdown of the Derby engine) and the 08006 error (successful shutdown of a single database) are the only exceptions thrown by Derby that might indicate that an operation succeeded. All other exceptions indicate that an operation failed. You should check the log file to be certain.

The errorPrint and SQLExceptionPrint methods

DERBY EXCEPTION REPORTING CLASSES: The two methods at the end of the file, `errorPrint` and `SQLExceptionPrint`, are generic exception-reporting methods that can be used with any JDBC program. This type of exception handling is required because often multiple exceptions (`SQLException`) are chained together and then thrown. A `while` loop is used to report on each error in the chain. The program starts this process by calling the `errorPrint` method from the `catch` block of the code that accesses the database.

```
//  Beginning of the primary catch block: uses errorPrint method
}  catch (Throwable e)  {
   /*    Catch all exceptions and pass them to
   **    the exception reporting method   */
   System.out.println(" . . . exception thrown:");
   errorPrint(e);
}
```

The `errorPrint` method prints a stack trace for all exceptions except a `SQLException`. Each `SQLException` is passed to the `SQLExceptionPrint` method.

```
static void errorPrint(Throwable e) {
   if (e instanceof SQLException)
      SQLExceptionPrint((SQLException)e);
   else {
      System.out.println("A non SQL error occured.");
      e.printStackTrace();
   }
}  // END errorPrint
```

The `SQLExceptionPrint` method iterates through each of the exceptions on the stack. For each error, the method displays the codes, message, and stacktrace.

```
//  Iterates through a stack of SQLExceptions
static void SQLExceptionPrint(SQLException sqle) {
   while (sqle != null) {
      System.out.println("\n---SQLException Caught---\n");
      System.out.println("SQLState:   " + (sqle).getSQLState());
      System.out.println("Severity: " + (sqle).getErrorCode());
      System.out.println("Message:  " + (sqle).getMessage());
      sqle.printStackTrace();
      sqle = sqle.getNextException();
   }
}  //  END SQLExceptionPrint
```

To see the output produced by this method, type a wish-list item with more than 32 characters, such as `I wish to see a Java program fail`.

Activity 4: Create and run a JDBC program using the client driver and Network Server

This activity demonstrates the ease with which a program that embeds Derby can be modified for a client/server implementation that uses the Derby Network Server.

This activity assumes you have performed the preceding activities and have a working directory called `DERBYTUTOR`, and have copies of the program files from the `$DERBY_HOME/demo/programs/workingwithderby/` directory. A basic knowledge of the `WwdEmbedded.java` program and experience starting and connecting to the Network Server are helpful. You will need to use a text editor to create the `WwdClient.java` program.

Note: As a convenience, the `workingwithderby` directory contains a program, `WwdClientExample.java`, which has already been edited in the appropriate manner. You can use this program directly as it is, or you can compare the contents of this program to the `WwdClient.java` program that you will construct in this activity, to see how the changes are made.
You will create a Derby client program, `WwdClient.java`, by changing a few lines of the `WwdEmbedded.java` program. You can run the client program in multiple command shells, allowing simultaneous update from two or more sources.

You use two command windows (Server-Shell and Client-Shell) in this activity. You use the Server-Shell to start the Derby Network Server and display Network Server messages. You use the Client-Shell to edit, compile and run the newly created `WwdClient.java` program. You set the `CLASSPATH` environment variable in the Client-Shell to support the client JDBC program.

1. Create the `WwdClient` program using the following steps:
   a. Open a command window (Client-Shell).
   b. Change to the `DERBYTUTOR` directory.

c. Make a copy of the `WwdEmbedded.java` program called `WwdClient.java`, as shown in the following table.

Table 18. Command to copy the program

Operating System	Command
UNIX (Korn Shell)	`cp WwdEmbedded.java WwdClient.java`
Windows	`copy WwdEmbedded.java WwdClient.java`

d. Open the `WwdClient.java` file in a text editor and update the class name to reflect the new file name:

```
Original declaration
     public class WwdEmbedded

New declaration
     public class WwdClient
```

e. Edit the *DEFINE VARIABLES SECTION* of the program so that the `driver` variable contains the name of the Derby client driver class and the `connectionURL` variable contains the hostname and port number of the Network Server.

```
Original definitions
     String driver = "org.apache.derby.jdbc.EmbeddedDriver";
     String dbName="jdbcDemoDB";
     String connectionURL = "jdbc:derby:" + dbName +
 ";create=true";

New definitions
     String driver = "org.apache.derby.jdbc.ClientDriver";
     ...
     String connectionURL = "jdbc:derby://localhost:1527/" +
 dbName + ";create=true";
```

f. Compile the application.

```
javac WwdClient.java
```

A command prompt appears if the compilation is successful. The binary file `WwdClient.class` is created. If an error message appears, modify the line indicated so that it is identical to the example. (As noted above, you can compare the contents of your `WwdClient.java` against the `WwdClientExample.java` program provided in the `workingwithderby` directory to verify that your editing was performed successfully.)

2. Set up the client/server environment using the following steps:
   a. Open a command window (Server-Shell).
   b. Change to the `DERBYTUTOR` directory.
   c. Start the Network Server, as shown in the following table.

Table 19. Command to start the Network Server

Operating System	Command

UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar server start`
	`Fri Sep 14 09:55:23 EDT 2012 : Security manager installed using the Basic server security policy.`
	`Fri Sep 14 09:55:24 EDT 2012 : Apache Derby`

Operating System	Command
	`Network Server - 10.10.0.0 - (1384314) started and ready to accept connections on port 1527`

Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar server start`
	`Fri Sep 14 09:55:23 EDT 2012 : Security manager installed using the Basic server security policy.`
	`Fri Sep 14 09:55:24 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) started and ready to accept connections on port 1527`

3. Run the client program using the following steps:
   a. Return to the Client-Shell window.
   b. Set the `CLASSPATH` environment variable to include the location of the file `derbyclient.jar` (not `derby.jar` as in the embedded example), as shown in the following table.

   Table 20. Command to set the `CLASSPATH` environment variable

Operating System	Command
UNIX (Korn Shell)	`export CLASSPATH=$DERBY_HOME/lib/derbyclient.jar:.`
Windows	`set CLASSPATH=%DERBY_HOME%\lib\derbyclient.jar;.`

   > Important: Include the dot (`.`) at the end of the command so that your current working directory is included in the classpath.

   c. Run the program:

```
java WwdClient
org.apache.derby.jdbc.ClientDriver loaded.
Connected to database jdbcDemoDB
Enter wish-list item (enter exit to end):
a sunny day
 ____________________________________________________
On 2012-09-14 09:46:37.77 I wished for a peppermint stick
On 2012-09-14 09:46:51.654 I wished for a long vacation
On 2012-09-14 09:56:36.6 I wished for a sunny day
 ____________________________________________________
Enter wish-list item (enter exit to end):
a new car
 ____________________________________________________
On 2012-09-14 09:46:37.77 I wished for a peppermint stick
On 2012-09-14 09:46:51.654 I wished for a long vacation
On 2012-09-14 09:56:36.6 I wished for a sunny day
On 2012-09-14 09:56:48.367 I wished for a new car
 ____________________________________________________
Enter wish-list item (enter exit to end):
exit
Closed connection
Getting Started With Derby JDBC program ending.
```

4. Shut down the Network Server, as shown in the following table.

   Table 21. Command to shut down the Network Server

Operating System	Command
UNIX (Korn Shell)	`java -jar $DERBY_HOME/lib/derbyrun.jar server shutdown` `Fri Sep 14 09:58:12 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) shutdown`
Windows	`java -jar %DERBY_HOME%\lib\derbyrun.jar server shutdown` `Fri Sep 14 09:58:12 EDT 2012 : Apache Derby Network Server - 10.10.0.0 - (1384314) shutdown`

The server shutdown confirmation appears in both command windows.

Activity notes

In a client/server environment, the client program is often used from other computers on the network. Whenever a system accepts connections from other computers, there is a chance of abuse. To maintain security, the Derby Network Server defaults to accepting connections only from clients running on the local machine (`localhost`). Before this or any other Derby client program can access the Network Server from another machine, additional steps should be taken to secure the Network Server environment. Once secured, the Network Server can be safely configured to accept connections from other machines. Refer to the "Network Server security" and "Running the Network Server under the security manager" sections of the *Java DB Server and Administration Guide* for important information on securing the Network Server and enabling network connections.

With the Network Server started, you can run the client program simultaneously in multiple windows. To demonstrate this, open two command windows and perform the substeps of the "Run the client program" step in each window. Both clients will operate without a problem. In contrast, it would not be possible for a program that uses the embedded driver (for example, `WwdEmbedded.java`) to access the database until the database or the Network Server is shut down.

You may have noticed that the client program does not shut down the database. This is because the database is a shared resource in a client/server environment and, in most cases, should be shut down only when the Network Server is shut down. If multiple clients are accessing the database and one client shuts down the database, the remaining clients will encounter a failure the next time they attempt an SQL command.

Derby's two architectures have caused confusion for some new Derby users, who mistakenly think that embedded is a single-user configuration. This is not true. The embedded driver supports multiple simultaneous connections, performs locking, and provides performance, integrity, and recoverability. Any application that uses the embedded driver can open multiple Derby connections and then provide a means for multiple users to interact with the database on each connection. The Derby Network Server is an example of such an application.

What next with Derby?

Congratulations on completing the activities in this tutorial. You now have experience with using Derby in both the embedded and client/server architectures.

With this basic knowledge, you are ready to begin using Derby to address your own specific needs. To learn more about this lightweight and powerful tool, refer to the Derby documentation, described in Product documentation for Java DB.

Also, take a few moments to become familiar with the many online resources available to Derby users and developers by browsing the Derby Web site at Apache as well as the

Derby Wiki. Whether you are performing a general evaluation of Derby or have a specific need to address, these links are a good stepping stone to finding additional information of interest:

- The "Learning material and other references" section on the front page of the Derby Wiki provides links to tutorials and other useful materials.
- The Derby Quick Start page is a good reference page organized by area of interest.
- You will find that many content-rich areas, such as the Derby Users mailing list, are available to you.
- If you are interested in how others are using Derby, see the Uses of Derby page on the Wiki. This page contains informational links to development projects and products that use Derby. When you implement a system using Derby, please add it to this list.

Activities summary

We hope these activities have helped you understand the steps needed to create and access Derby databases. Although Derby is easy to set up and use, you will find that it has the features and reliability of much larger database systems. The examples presented here do not begin to scratch the surface of what can be done.

Quick start guide for experienced JDBC users

This section is for experienced JDBC programmers who already know how to set the classpath, how to run a Java program, and how to use a JDBC driver.

For more detailed information on the topics covered in this section, see the *Java DB Developer's Guide* and the *Java DB Tools and Utilities Guide*. See Product documentation for Java DB for an overview of the Derby documentation.

Environments in which Derby can run

Before you configure your system to run Derby, it is useful to understand something about the different environments in which Derby can run, because these environments affect the classpath, driver name, and database connection URL.

See the *Java DB Developer's Guide* for more information on Derby environments.

Embedded environment

An embedded environment is an environment in which only a single application can access a database at one time, and no network access occurs.

When an application starts an instance of Derby within its JVM, the application runs in an embedded environment. Loading the embedded driver starts Derby.

Client/server environment

A client/server environment is an environment in which multiple applications connect to Derby over the network.

Derby runs embedded in a server framework that allows multiple network connections. (The framework itself starts an instance of Derby and runs in an embedded environment. However, the client applications do not run in the embedded environment.)

You can also embed Derby in any Java server framework.

See the *Java DB Server and Administration Guide* for more information on how to run Derby on a server.

Available drivers

Different JDBC drivers are available depending on the environment you choose for Derby.

- org.apache.derby.jdbc.EmbeddedDriver

  A driver for embedded environments, when Derby runs in the same JVM as the application.
- org.apache.derby.jdbc.ClientDriver

  A driver for the Network Server environment. The Network Server sets up a client/server environment.

Database connection URL

You must use a database connection URL when using the Derby-provided embedded driver to connect to a database.

The format for the database connection URL for connecting to a database is:

```
jdbc:derby:databaseName;URLAttributes
```

where:

- *databaseName*

  The name of the database that you want to connect to
- *URLAttributes*

  One or more of the supported attributes of the database connection URL, such as *;territory=ll_CC* or *;create=true*.

  For more information, see the *Java DB Developer's Guide*.

For the network client driver that is provided by Derby, the format for the database connection URL for connecting to a database is this:

```
jdbc:derby://server[:port]/databaseName[;URLAttributes=value[;...]]
```

where the *server* and *port* specify the host name (or IP address) and port number where the server is listening for requests and *databaseName* is the name of the database you want to connect to. The *URLAttributes* can be either Derby embedded or network client attributes. See the *Java DB Server and Administration Guide* for more information on accessing the Network Server by using the network client.

Documentation conventions

This section describes the terminology, syntax, and typographical conventions of the Derby documentation.

Terminology

The Derby documentation uses the specialized term `environment` to describe the method your application uses to interact with Derby.

An environment is sometimes referred to as a *framework*. The two types of environments are the *embedded environment* and the *client/server environment.*

SQL syntax

SQL syntax is presented in modified BNF notation.

The meta-symbols of BNF are listed in the following table.

Table 22. BNF meta-symbols

Symbol	Meaning
`|`	Or. Choose one of the items.
`[ ]`	Encloses optional items.
`*`	Flags items that you can repeat 0 or more times. Has a special meaning in some SQL statements.
`{ }`	Groups required items so that they can be marked with the symbol `|`.
`( ) . ,`	Other punctuation that is part of the syntax.

The following example shows how SQL syntax is presented:

```
CREATE [ UNIQUE ] INDEX indexName
  ON tableName ( simpleColumnName [ , simpleColumnName ]* )
```

Command-line syntax for running Java programs and utilities (as well as examples) always begins with the command *java*:

```
java org.apache.derby.tools.ij
```

This documentation attempts to be JVM neutral, but any command line examples or syntax using JVM specific references should be verified with your JVM documentation.

Typographical conventions

This documentation uses some typographical conventions to highlight elements of the SQL language, operating system commands, and the Java programming language.

The following table lists the typographical conventions used in the Derby documentation.

Table 23. Derby typographical conventions

Usage	Typeface	Examples
New terms	Italic	defined by *keys*
File and directory names	Italic	*C:\derby*
Dictionary objects	Italic	The *Employees* table
In syntax, items that you do not type exactly as they appear, but replace with the appropriate name	Italic	`CREATE TABLE tableName`
SQL examples	Bold and/or fixed-width	**`SELECT city_name FROM Cities`**
Java application examples	Bold and/or fixed-width	`Connection conn = DriverManager.getConnection ("jdbc:derby:Sample")`
Things you type in a command prompt	Bold and/or fixed-width	**`java org.apache.derby.tools.ij`**
Comments within examples	Bold and/or fixed-width	`--This line ignored`
SQL keywords (commands)	All caps	You can use a CREATE TABLE statement

Derby libraries and scripts: complete reference

This section describes Derby libraries and scripts.

Libraries provided by Derby

This section shows the different libraries used by Derby and their purposes.

The following table shows the Derby engine library.

Table 24. Engine library

Library Name	Use
`derby.jar`	For embedded databases. You always need this library for embedded environments. For client/server environments, you only need this library on the server.

For embedded environments, you need a library in the classpath to use a tool. For a client/server environment, you need a library on the client only. The following table shows the tools libraries.

Table 25. Tools libraries

Library Name	Use
`derbytools.jar`	Required for running all the Derby tools (such as `ij`, `dblook`, and import/export).
`derbyrun.jar`	Executable jar file that can be used to start the Derby tools.

The following table shows the Derby Network Server library.

Table 26. Network Server library

Library Name	Use
`derbynet.jar`	Required to start the Derby Network Server.

The following table shows the Derby network client library.

Table 27. Network client library

Library Name	Use
`derbyclient.jar`	Required to use the Derby network client driver.

The following table shows the Derby locale libraries.

Table 28. Locale libraries

Library Name	Use
derbyLocale_cs.jar	Required to provide translated messages for the Czech locale.
derbyLocale_de_DE.jar	Required to provide translated messages for the German locale.
derbyLocale_es.jar	Required to provide translated messages for the Spanish locale.
derbyLocale_fr.jar	Required to provide translated messages for the French locale.
derbyLocale_hu.jar	Required to provide translated messages for the Hungarian locale.
derbyLocale_it.jar	Required to provide translated messages for the Italian locale.
derbyLocale_ja_JP.jar	Required to provide translated messages for the Japanese locale.
derbyLocale_ko_KR.jar	Required to provide translated messages for the Korean locale.
derbyLocale_pl.jar	Required to provide translated messages for the Polish locale.
derbyLocale_pt_BR.jar	Required to provide translated messages for the Brazilian Portuguese locale.
derbyLocale_ru.jar	Required to provide translated messages for the Russian locale.
derbyLocale_zh_CN.jar	Required to provide translated messages for the Simplified Chinese locale.
derbyLocale_zh_TW.jar	Required to provide translated messages for the Traditional Chinese locale.

Scripts included with Derby

Derby includes scripts that start the Derby tools, add the appropriate libraries to the CLASSPATH, and start and stop the Network Server. These scripts available only when you download the *bin* distribution of Derby.

The Derby scripts are located in the `/bin` directory where you installed Derby. To run the scripts, you need to set several environment variables so that the scripts run correctly. See Setting the environment variables to set the environment variables.

You must run the Derby scripts in an appropriate command line environment, such as a Windows command prompt or a UNIX shell.

The complete list of scripts that are included with Derby is:

bin/dblook
: Runs the `dblook` tool.

bin/ij
: Starts the `ij` tool.

bin/NetworkServerControl
: Runs NetworkServerControl.

bin/setEmbeddedCP

Puts all of the Derby libraries for an embedded environment in the CLASSPATH.

bin/setNetworkClientCP

Puts the libraries needed to connect to the Derby Network Server into the CLASSPATH.

bin/setNetworkServerCP

Puts the libraries needed to start the Derby Network Server into the CLASSPATH.

bin/startNetworkServer

Starts the Network Server on the local machine.

bin/stopNetworkServer

Stops the Network Server on the local machine.

bin/sysinfo

Runs the `sysinfo` tool.

Trademarks

The following terms are trademarks or registered trademarks of other companies and have been used in at least one of the documents in the Apache Derby documentation library:

Cloudscape, DB2, DB2 Universal Database, DRDA, and IBM are trademarks of International Business Machines Corporation in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Other company, product, or service names may be trademarks or service marks of others.